

Theory Assignment 3

CS5280

Darpan Gaur
CO21BTECH11004

Problem 4.1

$$\begin{aligned}s_1 &= w_1(x)r_2(y)r_1(x)c_1r_2(x)w_2(y)c_2 \\ s_2 &= r_1(x)r_2(x)w_3(x)w_4(x)w_1(x)c_1w_2(x)c_2c_3c_4\end{aligned}$$

2PL

s1

$$s_1 = wl_1(x)w_1(x)rl_2(y)r_2(y)r_1(x)wu_1(x)rl_2(x)r_2(x)c_1wl_2(x)wl_2(y)w_2(y)wu_2(y)ru_2(y)c_2$$

s2

We need to abort one of transactions t_1 or t_2 . As there will be a deadlock condition, $rl_1(x)r_1(x)rl_2(x)r_2(x)$ after this point, t_1 and t_2 will be waiting for each other to release locks. So, we need to abort one of the transactions. Let's abort t_2 . So, output for s_2 is:

$$\begin{aligned}s_2 &= rl_1(x)r_1(x)rl_2(x)r_2(x)a_2wl_1(x)w_1(x)wu_1(x)ru_1(x)c_1 \\ &\quad wl_3(x)w_3(x)wu_3(x)wl_4(x)w_4(x)wu_4(x)c_3c_4\end{aligned}$$

O2PL

For s_1 we can use the same schedule as 2PL. For s_2 , we need to abort t_1 and t_2 , as t_3 and t_4 will share locks. So for operations $w_1(x)$ and $w_2(x)$, we need to wait until t_3 and t_4 release locks. But t_3 and t_4 will be waiting for t_1 and t_2 to release locks. So, we need to abort t_1 and t_2 . So, output for s_2 is:

$$s_2 = rl_1(x)r_1(x)rl_2(x)r_2(x)wl_3(x)w_3(x)wl_4(x)w_4(x)a_1a_2wu_3(x)wu_4(x)c_3c_4$$

BTO

s1

Timestamp of t_1 is less than t_2 . Also, we have only one conflicting operation $w_1(x)$ and $r_2(x)$ where $w_1(x) <_s r_2(x)$, and $ts(t_1) < ts(t_2)$. Hence, s_1 can be executed in BTO. Output schedule for s_1 is same as 2PL.

s2

Ordering of timestamps for t_1, t_2, t_3, t_4 is $t_1 < t_2 < t_3 < t_4$. Here, first $r_1(x), r_2(x), w_3(x)$ and $w_4(x)$ will be executed in order. Then for $w_1(x)$ we will check timestamps of t_1 and t_4 and since $ts(t_1) < ts(t_4)$, t_1 has to abort. Similarly, for $w_2(x)$ we will check timestamps of t_2 and t_4 and since $ts(t_2) < ts(t_4)$, t_2 has to abort. So output for s_2 :

$$s_2 = r_1(x)r_2(x)w_3(x)w_4(x)a_1a_2c_3c_4$$

SGT

s1

There is only one conflicting pair $w_1(x)$ and $r_2(x)$, so only one edge in the SGT graph. Hence, no need to abort any transaction. So output for s_1 is same as 2PL.

s2

In figure 1, we can see that there are cycle bwtween nodes $(t_1, t_3), (t_1, t_4), (t_2, t_3)$ and (t_2, t_4) . So t_1 and t_2 need to abort. So output for s_2 is:

$$s_2 = r_1(x)r_2(x)w_3(x)w_4(x)a_1a_2c_3c_4$$




Figure 1: Serial (Conflict) Graph for s_2

Problem 4.9

In figure 2, t_1 will acquire lock on y and t_2 will acquire lock on x . As we can share locks in O2PL, t_2 will be able to acquire lock on y and t_1 will be able to acquire lock on x . But while releasing locks, we need to preserve order of locks. Now, t_1 will be waiting to release lock on x by t_2 and t_2 will be waiting to release lock on y by t_1 . Hence, there will be a deadlock.




Figure 2: O2PL deadlock

Problem 4.12

The condition is seemingly natural but would lead to incorrect behaviour. If we remove a node, there is a possibility in future a cycle might be formed using the removed node, which will lead to incorrect behaviour of SGT protocol.

$$s = w_1(x)r_2(y)w_1(y)c_1r_3(z)w_2(z)c_2r_3(x)c_3$$

In figure 3, serialization graph is made using SGT protocol, and a cycle is formed. So, need to abort one of the transaction.

But if we use the given protocol in the question, we will remove t_1 after c_2 as t_1 is committed and active transaction at commit of t_1 , i.e., t_2 is also committed. So, we will remove t_1 and the serialization graph will look like figure 4, which is acyclic (dashed line shows removed nodes). Hence suggested protocol is incorrect.




Figure 3: Without removing nodes




Figure 4: With removing nodes

Problem 4.16

$$s = r_1(x)r_2(x)r_1(y)r_3(x)w_1(x)w_1(y)c_1r_2(y)r_3(z)w_3(z)c_3r_2(z)c_2$$

BOCC

In this case, t_2 and t_3 will be aborted.

- t_2 abort: Here, val/write phase of t_1 is before t_2 . $RS(t_2) \cap WS(t_1) \neq \phi$, as t_2 reads x and t_1 writes x . So, t_2 will be aborted.
- t_3 abort: Here, val/write phase of t_1 is before t_3 . $RS(t_3) \cap WS(t_1) \neq \phi$, as t_3 reads x and t_1 writes x . So, t_3 will be aborted.

Final schedule will be:

$$s = r_1(x)r_2(x)r_1(y)r_3(x)w_1(x)w_1(y)c_1r_2(y)r_3(z)a_3r_2(z)a_2$$

FOCC

In this case t_1 will be aborted.

- t_1 abort: As $WS(t_1) \cap RS(t_2) \neq \phi$, as t_1 writes x and t_2 reads x . Also, $WS(t_1) \cap RS(t_3) \neq \phi$, as t_1 writes x and t_3 reads x . So, t_1 will be aborted.

Final schedule will be:

$$s = r_1(x)r_2(x)r_1(y)r_3(x)a_1r_2(y)r_3(z)w_3(z)c_3r_2(z)c_2$$