

Hardware Architecture for Deep Learning

Homework 3: TinyML test on board

Darpan Gaur Lakshay Arora
CO21BTECH11004 CS24RESCH11006

Training tutorial

Setup the STM32CubeIDE, ST-Link, and the board and camera as per the tutorial. Created environment and installed the required packages. Followed the tutorial.

URAT input

- The instructions to send URAT input to the board is not mentioned in the tutorial.
- So, searched for it and found that it can be done using command shell console in STM32CUBEIDE.
- Figure 1 shows the steps to open command shell console.




Figure 1: URAT input

- Select connection name as Serial Port. (Figure 2).
- Click on New.. and enter connection name, select the serial port with which board is connected. (Figure 3).
- Finish and click OK.

- Enter "3" then "1" for training class 1.
- Enter "3" then "2" for training class 0.
- Enter "4" for inference.




Figure 2: URAT input

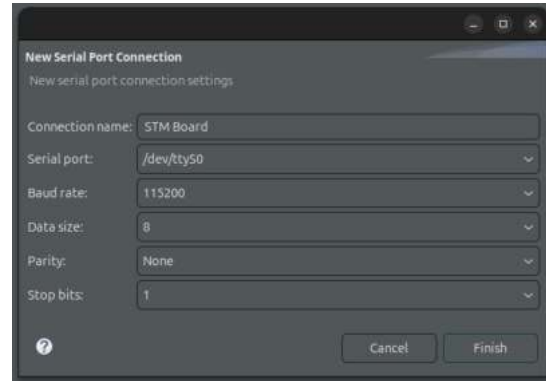


Figure 3: URAT input

Results Training

Figure 4 and 8 shows the training result of class 1 (person).
Figure 5, 9, and 7 shows the training result of class 0 (no person).




Figure 4: Training class 1




Figure 5: Training class 0



Figure 6: Training class 0




Figure 7: Training class 0




Figure 8: Training class 1

Inference tutorial

Results Inference

Figure 9 shows the inference result of a person without a camera.

Figure 10, 11 shows the inference result of a person with a camera.

Figure 12 shows the inference result of no person.




Figure 9: Person (Without camera)



Figure 10: Person




Figure 11: Person




Figure 12: No person

Errors Encountered

- STM32CUBEIDE installations gives error on Ubuntu 23.04 (missing library libncurses5.so). So, installed 24.04, on it works fine (But recommended is LTS 20.04 or 22.04).
- Need to install java and additional libraries for STM32CUBEIDE to work.
- Need to install ST-Link for communication between STM32 board and the host system, not mentioned in the tutorial.
- Tutorial mentioned to use python 3.6, but not able to install it so, used latest version 3.12.3.
- Getting overflow error while preparing the codebase by running python file for both training (tiny_training.py) and testing (vww.py), solved by changing data type of the variables in the code. Still some warnings were there while building the codebase.
- Not able to send URAT input to board (see training section).

Learnings

- Embedded Systems Development and Toolchain Setup
 - Gained hands-on experience working with the STM32CubeIDE, flashing firmware onto STM32 boards, and interfacing with external hardware like the Arducam camera.
 - Understood the importance of setting up the development environment properly — including installing required dependencies like Java, ST-Link utilities, and matching compatible OS versions (Ubuntu 24.04 worked better than 23.04, though the recommended was LTS 20.04/22.04).
- Challenges with Legacy Tools and Compatibility
 - Faced difficulties due to incompatibility of Python 3.6 with ARM-based processors and resolved it by using Python 3.12.3, after modifying the code to avoid overflow errors.
 - Learned that open-source repositories, especially older ones, often have missing dependencies, deprecated packages, or incomplete instructions, which require independent troubleshooting.
- Understanding UART and Serial Communication
 - Learned the functioning of UART (Universal Asynchronous Receiver/Transmitter) as a serial communication protocol and how to use the STM32CubeIDE's terminal to send commands to the board.
 - Realized the critical importance of matching the baud rate between transmitter and receiver to ensure reliable data transmission and avoid corruption.

- Hardware Constraints and Memory Management
 - Discovered that flashing our own code erased the factory-installed demo applications. This highlighted the limited storage capacity of the board and the importance of memory management in embedded systems.
 - Learned that real-time inference directly on the device results in significantly lower frames per second (FPS) compared to models trained separately and then deployed — due to computational limitations of microcontrollers.
 - While Training it resulted in 1.865 FPS, whereas while Inference it was 3.144 FPS.
- Hardware-Software Integration and Caution with Components
 - Successfully wired and interfaced the STM32 board with a camera, managing power and signal connections.
 - Realized how easy it is to misconnect wires, potentially damaging delicate components. This taught us the value of patience, precision, and consulting documentation or board schematics.
- Working with Edge AI and On-Device Learning
 - Understood the flow of on-device training vs. inference and the resource constraints (like RAM and CPU) that limit model complexity and speed on microcontrollers.
 - Got practical exposure to deploying AI models on hardware, and the trade-offs involved when choosing model size, accuracy, and inference time.