

# NSM GPU Programming

## Assignment 1

### Deadline : December 11, 2023, 23:55

## 1 Problem Statement

Write three separate CUDA C++ kernels for performing computations on two input matrices (**A** and **B**) and generating the output matrix **C**. In the first kernel **per\_row\_column\_kernel**, each thread should process a complete row of the input matrices. In the second kernel **per\_column\_row\_kernel**, each thread should process a complete column of the input matrices. In the third kernel **per\_element\_kernel**, each thread should process exactly one element from both the input matrices. For the evaluation purpose, **per\_row\_column\_kernel** will be invoked with *1D grid* and *1D blocks*, **per\_column\_row\_kernel** will be invoked with *1D grid and 2D blocks* and **per\_element\_kernel** will be invoked with *2D grid and 2D blocks*.

## 2 Input and Output

### 2.1 Input

- \* Matrix A of size m x n
- \* Matrix B of size m x n

### 2.2 Output

- \* Output is Matrix C of size n x m
- \* Output is computed as :  $C = (A + B)^T - (B - A)^T$ , where  $X^T$  is the transpose of matrix X.

### 2.3 Constraints

- \*  $2 \leq m \leq 2^{13}$  ,  $2 \leq n \leq 2^{13}$

## 3 Sample TestCase

- \* Input Matrix A

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- \* Input Matrix B

$$\begin{bmatrix} 1 & 3 & 2 \\ 4 & 7 & 2 \end{bmatrix}$$

- \*  $(A + B)^T$

$$\begin{bmatrix} 2 & 8 \\ 5 & 12 \\ 5 & 8 \end{bmatrix}$$

\*  $(B - A)^T$

$$\begin{bmatrix} 0 & 0 \\ 1 & 2 \\ -1 & -4 \end{bmatrix}$$

\*  $C = (A + B)^T - (B - A)^T$

$$C = \begin{bmatrix} 2 & 8 \\ 5 & 12 \\ 5 & 8 \end{bmatrix} - \begin{bmatrix} 0 & 0 \\ 1 & 2 \\ -1 & -4 \end{bmatrix}$$

$$= \begin{bmatrix} 2 & 8 \\ 4 & 10 \\ 6 & 12 \end{bmatrix}$$

\* Output Matrix C

$$\begin{bmatrix} 2 & 8 \\ 4 & 10 \\ 6 & 12 \end{bmatrix}$$

## 4 Points to be noted

- \* The number of threads launched for each of the three kernels will be more than or equal to the number of threads required to do the computation.
- \* Do not write any print statements inside the kernel.
- \* Test your code on large inputs.

## 5 Submission Guidelines

- \* Submit your file with your **full\_name.cu** which contains the implementation of the above-described functionality
- \* For transpose don't use tiling method.
- \* After submission, download the file and make sure it was the one you intended to submit.
- \* Kindly adhere strictly to the above guidelines.

## 6 Learning Suggestions

Write a CPU-version of code achieving the same functionality. Time the CPU code and GPU code separately for large matrices and compare the performances.