

NSM GPU Programming

Assignment 3

Deadline : December 13, 2023, 23:55

1 Problem Statement

Implement a parallel PageRank in CUDA. PageRank of a node p_i in a graph is given by the following formula.

$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

Figure 1: PR formula

Here p_1, p_2, \dots, p_N are the nodes, $M(p_i)$ is the set of nodes that link to p_i , $L(p_j)$ is the outdegree of node p_j , and N is the total number of nodes in the graph. $PR(p_j)$ should be the old value of pagerank of node p_j . Number of iterations of your pagerank implementation must be exactly 10 (i.e., the pagerank kernel should get invoked 100 times).

- * initialization of $PR(p_i) = 1/N$.
- * $d = 0.85$

2 Input and Output

2.1 Input

- * Input to your program would be a single argument mentioning path to the input graph in .txt format as given.
- * Input will be large graph datasets in .txt format

2.2 Output

- * Output of PageRank program should be N lines (N is the number of graph nodes) where i th line contains the PageRank value computed for the i th node. Print each PageRank value with nine decimal digits.
- * Output must be in .txt file

3 Points to be noted

- * We have given all the implementation except the CUDA part like kernel call, `cudamemcpy` etc. You need to write the corresponding code just below the commented section.
- * We have given a helper file **helper.hpp** which will keep track of CSR representation of the graph along with number of edges and number of nodes in a graph.

- * We have called CSR things like **offset array** and **neighbourlist array** in the main. You just need to pass it appropriately to kernel and complete the kernel part.
- * Do not write any print statements inside the kernel.
- * Don't write code having race condition.
- * You can use global synchronization also.
- * Do not upload anything other than the **full_name.zip** file. The zip file will consist of only helper.cu along with **your_name.cu** file.

4 Submission Guidelines

- * Submit your file with your **full_name.zip** which contains the implementation of the above-described functionality
- * Don't modify or change helper.hpp file
- * Don't use anything other than what is given in the .cu file Follow the naming conventions given for **CSR representation** of the graph.
- * Don't change the function prototype, function call, kernel argument and kernel call.
- * You only need to write the code just below the commented section
- * After submission, download the file and make sure it was the one you intended to submit.
- * Kindly adhere strictly to the above guidelines.

5 Learning Suggestions

- * Write a CPU-version of code achieving the same functionality. Time the CPU code and GPU code separately for large matrices and compare the performances.
- * Usage of global synchronisation